

STAT529
Applied Bayesian Decision Theory
Homework 4

Fangda Li
li1208@purdue.edu

April 28, 2017

1 Interim Analysis

Letting $\omega = \theta_T - \theta_S$, we aim to find the predictive probability of $P(\omega \geq 30) \geq 0.9$ at stage q in a total of n stages ($r = n - q$ denotes the rest of stages). According to the notes, the condition holds true when the following is satisfied:

$$\begin{aligned}\Delta_n &= \frac{q}{n}\Delta_q + \frac{r}{n}\Delta_r \geq 30 + 1.3\sigma\sqrt{\frac{2}{n}}, \\ \Delta_r &\geq \frac{n}{r}(30 + 1.3\sigma\sqrt{\frac{2}{n}} - \frac{q}{n}\Delta_q).\end{aligned}$$

Taking $q = r = 100$, $n = 200$ and $\sigma = 20$, the above condition is equivalent to:

$$\Delta_r \geq 65.2 - \Delta_q.$$

Also recall that the posterior distribution of ω can be shown to be:

$$\pi(\omega|\Delta_q, \beta, \tau^2) \sim N(m, v^2),$$

where $m = \frac{\frac{2\sigma^2}{q}(0) + 2\tau^2\Delta_q}{\frac{2\sigma^2}{q} + 2\tau^2}$ and $v^2 = \frac{\frac{2\sigma^2}{q}2\tau^2}{\frac{2\sigma^2}{q} + 2\tau^2}$. Then the posterior distribution of Δ_r with β and ω being integrated out becomes:

$$\begin{aligned}f(\Delta_r|\Delta_q) &= \int \int \int_{\tau^2, \beta, \omega} \frac{f(\Delta_r, \Delta_q|\omega, \beta, \tau^2)\pi(\omega|\beta, \tau^2)h_1(\beta)h_2(\tau^2)}{f(\Delta_q)} \\ &= \int \int \int_{\tau^2, \beta, \omega} f(\Delta_r, \omega)\pi(\omega|\Delta_q, \beta, \tau^2)f(\Delta_q|\beta, \tau^2)\frac{h_1(\beta)h_2(\tau^2)}{f(\Delta_q)} \\ &= \int_{\tau^2} N(m, \frac{2\sigma^2}{r} + v^2)\frac{f(\Delta_q|\tau^2)h_2(\tau^2)}{f(\Delta_q)}.\end{aligned}$$

With the posterior distribution, we can compute the posterior probability of interest:

$$P(\Delta_r \geq 65.2 - \Delta_q|\Delta_q) = \int_{\tau^2} (1 - \Phi(\frac{(65.2 - \Delta_q) - \Delta_q\frac{\tau^2}{4+\tau^2}}{\sqrt{8 + \frac{8\tau^2}{4+\tau^2}}}))\frac{f(\Delta_q|\tau^2)h_2(\tau^2)}{f(\Delta_q)}.$$

In practice, MCMC is performed in order to generate posterior samples of τ^2 and the predictive probability is computed by averaging the $(1 - \Phi(*))$ values. The probability r used in MCMC for keeping τ_t^2 is:

$$r = \min\{\frac{f(\Delta_q|\tau_t^2)}{f(\Delta_q|\tau_{t-1}^2)}, 1\}.$$

According to my simulation, the computed predictive probabilities $P(\Delta_r \geq 65.2 - \Delta_q)$ for different Δ can be found in Table 1:

Table 1: Predictive probabilities $P(\Delta_r \geq 65.2 - \Delta_q)$

Δ_q	30	32	34	36	38	40
P	0.008	0.078	0.325	0.699	0.932	0.993

2 Sequential Analysis

The problem is setup as follows: observations $X_1, X_2, \dots, X_n \sim \text{uniform}(0, 50)$, “payoff” function $Y_n = X_n - nc$, and the number of possible trials $N = 15$.

- a) **Backward Induction:** We start by looking at the second last stage, $n = 14$. First, the expected payoff at stage $n = 15$ is:

$$E[Y_{15}] = E[X_{15} - 15c] = 25 - 15c.$$

We will stop at stage 14 if the payoff at stage 14, y_{14} , is greater than $E[Y_{15}]$, which is $x_{14} \geq 25 - c$. Now we can define the optimal payoff at stage 14 as the larger value between the observed payoff y_{14} and the expected payoff at the next stage $E[Y_{15}]$: $\beta_{14}^{15} = \max(y_{14}, E[Y_{15}])$.

The optimal rule says we should stop at stage $n = 13$ if the observed payoff at stage 13, y_{13} , is greater than the expected optimal payoff at stage 14, which is:

$$\begin{aligned} y_{13} &\geq E[\beta_{14}^{15}] \\ x_{13} - 13c &\geq \int_0^{50} \max(y_{14}, E[Y_{15}]) f(x_{14}) dx_{14} \\ x_{13} - 13c &\geq \int_0^{50} \max(x_{14} - 14c, 25 - 15c) \frac{1}{50} dx_{14} \\ x_{13} &\geq \int_0^{50} \max(x_{14} - 14c, 25 - 15c) \frac{1}{50} dx_{14} + 13c \end{aligned}$$

Continuing on to stage 12, 11, ..., 1 using backward induction, we can therefore obtain the optimal stopping rules at each stage. For example, at stage 12, the optimal rule says we should continue as long as:

$$y_{12} \geq E[\beta_{13}^{15}] = E[\max(Y_{13}, E[\beta_{14}^{15}])].$$

In my simulation, I have chosen $c = 5$ for which it is worthwhile to take observations. According to the optimal stopping rule, we will stop at stage n if x_n is greater than the threshold t_n , where t_n for $n = 0, 1, \dots, 14$. The obtained t_n values for $c = 5$ are plotted in Figure 1, which also includes results for other choices of c for the sake of comparison. Note that $c = 25$, for example, provides us a scenario where no observations are taken. At stage 1, the rule will tell us to stop if any positive payoff (which is always the case).

- b) The expected payoff for the optimal stopping rule can be found in Figure 1. For $c = 5$, the expected payoff is $E[\beta_1^{15}] - c = \mathbf{22.633}$.




Figure 1: Optimal stopping thresholds for various choices of c .




Figure 2: Expected payoffs for optimal stopping rule and OSA.

c) **One Step Ahead (OSA)**: we stop at stage n where the following condition is met:

$$y_n \geq E[Y_{n+1}],$$

which expands to

$$x_n - nc \geq E[X_{n+1}] - (n+1)c$$

$$x_n - nc \geq 25 - (n+1)c$$

$$x_n \geq 25 - c.$$

Simply put, we stop as soon as the observation x_n is greater than $25 - c$. The expected payoff of the OSA stopping rule is obtained by “playing the game” many times and averaging the outcomes. The resulting OSA expected payoffs of different c can be also found in Figure 2. For $c = 5$, the expected payoff is **31.327**.

3 Secretary Problem

First of all, the following constants are provided: $\sigma^2 = 50$, $\tau^2 = 25$ and $\beta = 150$. In this problem, the observed performance of candidate n is $x_n \sim N(\theta_n | \sigma^2)$, where $\theta \sim N(\beta, \tau^2)$. Our goal is to hire the interviewee with the highest posterior mean:

$$\begin{aligned}
 E[\pi(\theta_n | x_n)] &= E\left[N\left(\frac{\sigma^2\beta + \tau^2x_n}{\sigma^2 + \tau^2}, \frac{\sigma^2\tau^2}{\sigma^2 + \tau^2}\right)\right] \\
 &= \frac{\sigma^2\beta + \tau^2x_n}{\sigma^2 + \tau^2} \\
 &= 100 + \frac{x_n}{3} \\
 &= w_n \sim N(150, 8.3).
 \end{aligned}$$

In other words, our payoff function is $Y_n = \max(w_1, w_2, \dots, w_n) - a_n$. Since the cost a_n only increments by a constant value c , our model is a Monotone Case, where the OSA stopping rule is optimal. Recall that OSA says we should stop when:

$$y_n \geq E[Y_{n+1}].$$

Letting m_n denote $\max(w_1, w_2, \dots, w_n)$, we can expand the condition to:

$$\begin{aligned} m_n - a_n &\geq E[m_{n+1} - a_{n+1}] \\ m_n &\geq E[m_{n+1}] - c \\ m_n &\geq m_n + E[(w_{n+1} - m_n)^+] - c \\ E[(w_{n+1} - m_n)^+] &\geq c. \end{aligned}$$

Now define γ such that $E[(w_{n+1} - \gamma)^+] = c$. Then the OSA stopping rule becomes: stop the first time $m_n \geq \gamma$, and we can solve for γ using the following integral equation:

$$\begin{aligned} \int_{\gamma}^{\infty} (w - \gamma) f(w) dw &= c \\ \int_{\gamma}^{\infty} (w - \gamma) \frac{1}{2.88\sqrt{2\pi}} e^{-\frac{1}{2} \frac{(w-150)^2}{8.3}} dw &= c. \end{aligned}$$

Using simulation, we arrive at $\gamma \approx 145$ when $c = 5$.

4 Code

4.1 Q1

```
from pylab import *
from scipy.stats import norm as normal

def get_tau2_posterior_samples_MCMC(deltaq):
    '''
        Get the posterior samples for tau2.
    '''
    N = 20000
    # Posteriors before selection
    tau2 = zeros(N)
    # Initial guesses
    tau2[0] = rand() * 20 + 10
    # Generating loop
    i = 1
    while i < N:
        # Generate new sample
        tau2_i = rand() * 20 + 10
        # Determine whether to keep
```

```

    numer = exp(-deltaq**2 / 2 / (8 + 2 * tau2_i)) / sqrt(8 + 2 * tau2_i)
    denom = exp(-deltaq**2 / 2 / (8 + 2 * tau2[i-1])) / sqrt(8 + 2 *
        ↪ tau2[i-1])
    r = min(numer / denom, 1)
    # Keep with probability r
    if rand() > r: continue
    tau2[i] = tau2_i
    i += 1
    # Discard the first 500 to minimize the influence of the initial guess
    tau2 = tau2[500:]
    return tau2

def main():
    set_printoptions(precision=3)
    P = zeros(6)
    for i, deltaq in enumerate([30, 32, 34, 36, 38, 40]):
        tau2 = get_tau2_posterior_samples_MCMC(deltaq)
        rv = normal()
        numer = 65.2 - deltaq - deltaq * tau2 / (4 + tau2)
        denom = sqrt(8 + 8 * tau2 / (4 + tau2))
        Ps = 1 - rv.cdf(numer / denom)
        P[i] = mean(Ps)
    print P

if __name__ == '__main__':
    main()

```

4.2 Q2

```

from pylab import *

def get_optimal_rules_backward_induction(c):
    '''
        Obtain the optimal stopping thresholds for each stage.
        Thresholds are designated for x.
    '''
    # Optimal thresholds from stage 0 to 14
    # Stage 15 doesn't have a rule
    thresh = zeros(15)
    thresh[14] = 25 - c # Stage 14
    E_beta = zeros(15)
    E_beta[14] = 25 - 15 * c # beta_14^15
    for n in range(13, 0, -1):
        thresh[n] = E_beta[n + 1] + n * c
        E_beta[n] = (E_beta[n + 1] + n * c)**2 / 100. + 25 - n * c

```

```

    return thresh[1:], E_beta[1] - c

def get_rules_osa(c):
    """
    Get OSA  $x$  thresholds, but no optimality guaranteed.
    """
    return ones(15) * (25 - c)

def get_expected_payoff_osa(c, thresh):
    """
    Repeat the game and return the average payoff.
    """
    N = 1000
    Y = zeros(N)
    for i in range(N):
        j = 0
        x = rand() * 50
        while j < 15 and x < thresh[j]:
            j += 1
        x = rand() * 50
        Y[i] = x - j * c
    return mean(Y)

def main():
    set_printoptions(precision=3, suppress=True)
    C = [0, 1, 2, 5, 10, 25, 50]
    Y_opt = zeros(len(C))
    Y_osa = zeros((len(C)))
    figure()
    for i, c in enumerate(C):
        t_opt, Y_opt[i] = get_optimal_rules_backward_induction(c)
        t_osa = get_rules_osa(c)
        Y_osa[i] = get_expected_payoff_osa(c, t_osa)
        print "c =", c
        print t_opt
        print "Optimal expected payoff =", Y_opt[i]
        print "OSA expected payoff =", Y_osa[i]
        step(arange(14) + 1, t_opt, label="c = %d" % c)
    title('Optimal Stopping Thresholds for  $x_i$ ')
    xlabel('Stage')
    ylabel('Threshold')
    legend()
    # Plot the expected payoffs
    figure()
    bar_width = 0.35

```



```
index = arange(len(C))
bar(index, Y_opt, bar_width, label='OPT')
bar(index + bar_width, Y_osa, bar_width, label='OSA')
title('Expected Payoffs: Optimal v.s. OSA')
xlabel('c')
ylabel('Payoff')
xticks(index, [str(c) for c in C])
legend()
tight_layout()
show()

if __name__ == '__main__':
    main()
```